

Kubernetes Linux Academy

Docker installation:

We are installing it on 3 nodes.

1 master:

bakumar2c.mylabserver.com

2 worker nodes:

bakumar3c.mylabserver.com,

bakumar4c.mylabserver.com

Run below commands in all 3 nodes

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
```

```
add-apt-repository \
"deb [arch=amd64] https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) \
stable"
```

```
apt-get update
```

```
apt-get install -y docker-ce=18.06.1~ce~3-0~ubuntu
```

```
apt-mark hold docker-ce
```

```
docker ps
```

The first step in setting up a new cluster is to install a container runtime such as Docker. In this lesson, we will be installing Docker on our three servers in preparation for standing up a Kubernetes cluster. After completing this lesson, you should have three playground servers, all with Docker up and running.

Here are the commands used in this lesson:

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -

sudo add-apt-repository \
"deb [arch=amd64] https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) \
stable"

sudo apt-get update

sudo apt-get install -y docker-ce=18.06.1~ce~3-0~ubuntu

sudo apt-mark hold docker-ce
```

You can verify that docker is working by running this command:

```
sudo docker version
```

[Introduction](#)[Building a Cluster](#)[Cluster Architecture](#)[Setting Up Playground Servers](#)[Installing Docker](#)[Installing Kubeadm, Kubelet, and Kubectl](#)[Bootstrapping the Cluster](#)[Configuring Networking with Flannel](#)[Basic Kubernetes Concepts](#)[Deploying to Kubernetes](#)[Microservices](#)[Additional Resources](#)

Cluster Architecture

In this section, I am going to show you how to build your own Kubernetes cluster.

There is a learning activity at the end of this section that you can use to build your cluster, but you can also follow along using playground servers if you want to.

Here is the architecture of the cluster that we will be building:

Kube Master

Docker

Kubeadm

Kubelet

Kubectl

Control Plane

Kube Node 1

Docker

Kubeadm

Kubelet

Kubectl

Kube Node 2

Docker

Kubeadm

Kubelet

Kubectl



Introduction

Building a Cluster

Cluster Architecture

Setting Up Playground
Servers

Installing Docker

Installing Kubeadm, Kubelet,
and Kubectl

Bootstrapping the Cluster

Configuring Networking with
Flannel

Basic Kubernetes Concepts

Deploying to Kubernetes

Microservices

Additional Resources

Setting Up Playground Servers

You can simply use the learning activity at the end of this section to practice setting up your own Kubernetes cluster, but you can also use the Linux Academy playground servers.

If you want to do that, create three playground servers with the following settings:

Distribution:

Ubuntu 18.04 Bionic Beaver LTS

Size: Small: 2 unit(s)

Tag:

For Server 1: Kube Master

For Server 2: Kube Node 1

For Server 3: Kube Node 2

Introduction

Building a Cluster

Cluster Architecture

Setting Up Playground Servers

Installing Docker

Installing Kubeadm, Kubelet, and Kubectl

Bootstrapping the Cluster

Configuring Networking with Flannel

Basic Kubernetes Concepts

Deploying to Kubernetes

Microservices

Additional Resources

Installing Kubeadm, Kubelet, and Kubectl

Next>

Now we are ready to install the components of Kubernetes itself.

We will be installing three things on each of our servers:

- **Kubeadm:** This is a tool which automates a large portion of the process of setting up a cluster. It will make our job much easier!
- **Kubelet:** The essential component of Kubernetes that handles running containers on a node. Every server that will be running containers needs kubelet.
- **Kubectl:** Command-line tool for interacting with the cluster once it is up. We will use this to manage the cluster.

Install below Kubernetes steps In all 3 nodes

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
```

```
cat << EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF
```

```
apt-get update
```

```
apt-get install -y kubelet=1.12.7-00 kubeadm=1.12.7-00 kubectl=1.12.7-00
```

```
apt-mark hold kubelet kubeadm kubectl
```

```
kubeadm version
```

Now that Docker is installed, we are ready to install the Kubernetes components. In this lesson, I will guide you through the process of installing Kubeadm, Kubelet, and Kubectl on all three playground servers. After completing this lesson, you should be ready for the next step, which is to bootstrap the cluster.

Here are the commands used to install the Kubernetes components in this lesson. Run these on all three servers.

NOTE: There are some issues being reported when installing version 1.12.2-00 from the Kubernetes ubuntu repositories. You can work around this by using version 1.12.7-00 for kubelet, kubeadm, and kubectl.

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -  
  
cat << EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list  
deb https://apt.kubernetes.io/ kubernetes-xenial main  
EOF  
  
sudo apt-get update  
  
sudo apt-get install -y kubelet=1.12.7-00 kubeadm=1.12.7-00 kubectl=1.12.7-00  
  
sudo apt-mark hold kubelet kubeadm kubectl
```

After installing these components, verify that Kubeadm is working by getting the version info.

```
kubeadm version
```

Bootstrapping Kubernetes cluster

Run Below commands in Master node.

```
kubeadm init --pod-network-cidr=10.244.0.0/16
```

<It will display set of further commands to execute on same master node along with token key to join worker nodes into cluster.>

Example:

[

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You can now join any number of machines by running the following on each node as root:

```
kubeadm join 172.31.41.71:6443 --token ocf5ib.xzicjfjw7slt271 --discovery-token-ca-cert-hash sha256:e6332502487041e1314afec51b7a31486edc12d1a41861c1380d8791816bf117  
]
```

Run above shown commands in master node:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Then for checking kubectl

```
#kubectl version
```

output:

```
Client Version: ...  
Server Version:...
```

Run above master node shown output of “token join” in worker node

From bakumar3c.mylabserver.com,

```
sudo kubeadm join 172.31.41.71:6443 --token ocf5ib.xzicjfw7slt271 --discovery-token-ca-cert-hash sha256:e6332502487041e1314afec51b7a31486edc12d1a41861c1380d8791816bf117
```

From bakumar4c.mylabserver.com

```
sudo kubeadm join 172.31.41.71:6443 --token ocf5ib.xzicjfw7slt271 --discovery-token-ca-cert-hash sha256:e6332502487041e1314afec51b7a31486edc12d1a41861c1380d8791816bf117
```

Once ran above commands, we run below command from Master node to see list of joined nodes

```
root@bakumar2c:~# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
bakumar2c.mylabserver.com	NotReady	master	8m5s	v1.12.7
bakumar3c.mylabserver.com	NotReady	<none>	3m46s	v1.12.7
bakumar4c.mylabserver.com	NotReady	<none>	3m31s	v1.12.7

Till we setup **networking** for cluster, “STATUS” will show up in “Not Ready”

Now we are ready to get a real Kubernetes cluster up and running! In this lesson, we will bootstrap the cluster on the KubeMaster node. Then, we will join each of the two worker nodes to the cluster, forming an actual multi-node Kubernetes cluster.

Here are the commands used in this lesson:

- On the KubeMaster node, initialize the cluster:

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

That command may take a few minutes to complete.

- When it is done, set up the local kubeconfig

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Verify that the cluster is responsive and that Kubectl is working

```
kubectl version
```

You should get **Server Version** as well as **Client Version**. It should look something like this:

```
Client Version: version.Info{Major:"1", Minor:"12", GitVersion:"v1.12.2", GitCommit:"17c77c78821897514c8e573582a8d2313dc74b", GitTreeState:"clean", BuildDate:"2018-10-24T04:54:59Z", GoVersion:"go1.10.4", Compiler:"gc", Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"12", GitVersion:"v1.12.2", GitCommit:"17c77c78821897514c8e573582a8d2313dc74b", GitTreeState:"clean", BuildDate:"2018-10-24T04:54:59Z", GoVersion:"go1.10.4", Compiler:"gc", Platform:"linux/amd64"}
```

- The **kubeadm init** command should output a **kubeadm join** command containing a token and hash. Copy that command and run it with **sudo** on both worker nodes. It should look something like this:

```
sudo kubeadm join $some_ip:6443 --token $some_token --discovery-token-ca-cert-hash $some_hash
```

- Verify that all nodes have successfully joined the cluster:

```
kubectl get nodes
```

You should see all three of your nodes listed. It should look something like this:

NAME	STATUS	ROLES	AGE	VERSION
vboxd1c.mylabserver.com	NotReady	master	5m17s	v1.12.2
vboxd2c.mylabserver.com	NotReady	<none>	51s	v1.12.2
vboxd3c.mylabserver.com	NotReady	<none>	31s	v1.12.2

Note: The nodes are expected to have a STATUS of **NotReady** at this point.

Kubernetes- Setting up networking with Flannel

Run below commands in all nodes-master & workers in cluster

```

root@bakumar2c:~# echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/
sysctl.conf
net.bridge.bridge-nf-call-iptables=1
root@bakumar2c:~# sysctl -p
net.bridge.bridge-nf-call-iptables = 1
root@bakumar2c:~#

```

Install Flannel in the cluster by running this only on the Master node:

```

root@bakumar2c:~# kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/
bc79dd1505b0c8681ece4de4c0d86c5cd2643275/Documentation/kube-flannel.yml

```

```

clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.extensions/kube-flannel-ds-amd64 created
daemonset.extensions/kube-flannel-ds-arm64 created
daemonset.extensions/kube-flannel-ds-arm created
daemonset.extensions/kube-flannel-ds-ppc64le created
daemonset.extensions/kube-flannel-ds-s390x created

```

Now we setted up networking with flannel, check the nodes status.

```

root@bakumar2c:~# kubectl get nodes

```

NAME	STATUS	ROLES	AGE	VERSION
bakumar2c.mylabserver.com	Ready	master	16m	v1.12.7
bakumar3c.mylabserver.com	Ready	<none>	12m	v1.12.7
bakumar4c.mylabserver.com	Ready	<none>	12m	v1.12.7

```

root@bakumar2c:~#

```

Check the namespace of “kube-system”, which by default all containers are created under this.

```

root@bakumar2c:~# kubectl get pods -n kube-system

```

NAME	READY	STATUS	RESTARTS	AGE
coredns-bb49df795-b4hzd	1/1	Running	0	19m
coredns-bb49df795-fgj4j	1/1	Running	0	19m
etcd-bakumar2c.mylabserver.com	1/1	Running	0	18m
kube-apiserver-bakumar2c.mylabserver.com	1/1	Running	0	18m
kube-controller-manager-bakumar2c.mylabserver.com	1/1	Running	0	18m
kube-flannel-ds-amd64-jmdxn	1/1	Running	0	3m9s
kube-flannel-ds-amd64-rn6xj	1/1	Running	0	3m9s
kube-flannel-ds-amd64-smn2n	1/1	Running	0	3m9s
kube-proxy-hxxbk	1/1	Running	0	19m
kube-proxy-l4n6s	1/1	Running	0	14m
kube-proxy-sqrkh	1/1	Running	0	15m
kube-scheduler-bakumar2c.mylabserver.com	1/1	Running	0	18m

The kube-flannel is responsible for setting up networking layer in kebernetes.

Once the Kubernetes cluster is set up, we still need to configure cluster networking in order to make the cluster fully functional. In this lesson, we will walk through the process of configuring a cluster network using Flannel. You can find more information on Flannel at the official site:

<https://www.kubernetes.io/en/docs/tutorials/network-configuration/flannel/>

Here are the commands used in this lesson:

- On all three nodes, run the following:

```
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
```

- Install Flannel in the cluster by running this only on the Master node:

```
kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/bc7de39c1b7f8a0d9c4515d28533790302275/Documentation/kube-flannel.yml
```

- Verify that all three nodes now have a STATUS of **Ready**:

```
kubectl get nodes
```

You should see all three of your servers listed, and all should have a STATUS of **Ready**. It should look something like this:

NAME	STATUS	ROLES	AGE	VERSION
vmay22eu1.labs.server.com	Ready	master	5m17s	v1.12.2
vmay22eu2.labs.server.com	Ready	worker	55s	v1.12.2
vmay22eu3.labs.server.com	Ready	worker	51s	v1.12.2

Note: It may take a few moments for all nodes to enter the **Ready** status, so if they are not all **Ready**, wait a few moments and try again.

- It is also a good idea to verify that the Flannel pods are up and running. Run this command to get a list of system pods:

```
kubectl get pods -n kube-system
```

You should have three pods with **flannel** in the name, and all three should have a status of **Running**.

Containers and Pods

Introduction

Building a Cluster

Basic Kubernetes Concepts

Containers and Pods

Clustering and Nodes

Networking in Kubernetes

Kubernetes Architecture and Components

Deploying to Kubernetes

Microservices

Additional Resources

Containers and Pods

Next>

Pods are the smallest and most basic building block of the Kubernetes model.

A **pod** consists of one or more containers, storage resources, and a unique IP address in the Kubernetes cluster network.

Node

Pod
10.244.0.1

Container

Pod
10.244.0.2

Container

Container

Node

Pod
10.244.0.3

Container

In order to run containers, Kubernetes schedules pods to run on servers in the cluster. When a pod is scheduled, the server will run the containers that are part of that pod.

[Sample YAML file to create nginx image](#)

Introduction

Building a Cluster

Basic Kubernetes Concepts

Containers and Pods

Clustering and Nodes

Networking in Kubernetes

Kubernetes Architecture
and Components

Deploying to Kubernetes

Microservices

Additional Resources

◀ Prev

Containers and Pods

Let's create a simple pod that runs an Nginx web server!

```
cat << EOF | kubectl create -f -  
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
spec:  
  containers:  
  - name: nginx  
    image: nginx  
EOF
```

```
kubectl get pods --all-namespaces
```

```
kubectl describe pod $pod_name -n $namespace
```



Run below in Master node to create a nginx image

```
cat << EOF | kubectl create -f -
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx
```

```
spec:
```

```
  containers:
```

```
    - name: nginx
```

```
      image: nginx
```

```
EOF
```

```
root@bakumar2c:/opt# kubectl get pods
```

```
NAME    READY   STATUS    RESTARTS   AGE
nginx   1/1     Running   0           24s
```

more info about created pod "nginx"

```
root@bakumar2c:/opt# kubectl describe pod nginx
```

```
Name:          nginx
Namespace:     default
Priority:       0
PriorityClassName: <none>
Node:          bakumar4c.mylabserver.com/172.31.47.144
Start Time:    Mon, 14 Oct 2019 06:46:27 +0000
Labels:        <none>
Annotations:    <none>
Status:        Running
IP:            10.244.2.2
Containers:
  nginx:
    Container ID:  docker://13c158c33b40ac51c3091b6e76662984c214d6ab5caf3bcd040a44431b9039a0
    Image:         nginx
    Image ID:      docker-pullable://nginx@sha256:aeded0f2a861747f43a01cf1018cf9efe2bdd02afd57d2b11fcc7cad616ccd1
    Port:          <none>
    Host Port:     <none>
    State:         Running
      Started:     Mon, 14 Oct 2019 06:46:38 +0000
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-chmls (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready          True
  ContainersReady True
  PodScheduled   True
Volumes:
  default-token-chmls:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-chmls
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type Reason Age From Message
  ----
Normal Scheduled 4m43s default-scheduler Successfully assigned default/nginx to bakumar4c.mylabserver.com
Normal Pulling   4m42s kubelet, bakumar4c.mylabserver.com pulling image "nginx"
Normal Pulled    4m34s kubelet, bakumar4c.mylabserver.com Successfully pulled image "nginx"
Normal Created   4m32s kubelet, bakumar4c.mylabserver.com Created container
Normal Started   4m32s kubelet, bakumar4c.mylabserver.com Started container
root@bakumar2c:/opt#
```

- [illegible]

- [vissuti per prima](#)

- KUMHETI, DEBENDRA BHAI MALIK

- Author's address: Department of Mathematics, University of Illinois at Chicago, Chicago, IL 60607, USA

```
root@bakumar2c:/opt#  
root@bakumar2c:/opt# kubectl delete pod nginx  
pod "nginx" deleted  
root@bakumar2c:/opt#
```

Clustering and nodes



Introduction

Building a Cluster

Basic Kubernetes Concepts

Containers and Pods

Clustering and Nodes

Networking in Kubernetes

Kubernetes Architecture and Components

Deploying to Kubernetes

Microservices

Additional Resources

Clustering and Nodes

Kubernetes implements a clustered architecture. In a typical production environment, you will have multiple servers that are able to run your workloads (containers).

These servers which actually run the containers are called nodes.

A Kubernetes cluster has one or more control servers which manage and control the cluster and host the Kubernetes API. These control servers are usually separate from worker nodes, which run applications within the cluster.



You can use `kubectl` to list nodes in the cluster and get more information about them:

```
kubectl get nodes
```

```
kubectl describe node $node_name
```

Master node -> Kubernetes Control nodes

Worker nodes are responsible for running actual applications

Nodes are an essential part of the Kubernetes cluster. They are the machines where your cluster's container workloads are executed. In this lesson, we will discuss node roles and labels, and how to explore more ways in which you can configure nodes about roles in your cluster.

- Details of nodes

```
kubectl get nodes
```

- Get more information about a specific node

```
kubectl describe node $node_name
```

Networking in Kubernetes

Introduction

Building a Cluster

Basic Kubernetes Concepts

Containers and Pods

Clustering and Nodes

Networking in Kubernetes

Kubernetes Architecture and Components

Deploying to Kubernetes

Microservices

Additional Resources

Networking in Kubernetes

When using Kubernetes, it is important to understand how Kubernetes implements networking between pods (and services) in the cluster.

The Kubernetes networking model involves creating a **virtual network** across the whole cluster. This means that every pod on the cluster has a unique IP address, and can communicate with any other pod in the cluster, even if that other pod is running on a different node.

Node	Node
Virtual Cluster Network	
Pod	Pod
Pod	Pod
	Pod

Kubernetes supports a variety of networking plugins that implement this model in various ways. In this course, we will be using **Flannel**.

Run it in Master node

```
cat << EOF | kubectl create -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.15.4
        ports:
        - containerPort: 80
EOF
```

root@bakumar2c:/opt# kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
nginx-d55b94fd-r25rc	1/1	Running	0	73s
nginx-d55b94fd-w2rmv	1/1	Running	0	73s


```
cat << EOF | kubectl create -f -
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: busybox
```

```
spec:
```

```
  containers:
```

```
  - name: busybox
```

```
    image: radial/busyboxplus:curl
```

```
    args:
```

```
    - sleep
```

```
    - "1000"
```

```
EOF
```

```
--
```

```
root@bakumar2c:/opt# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
busybox	1/1	Running	0	37s
nginx-d55b94fd-r25rc	1/1	Running	0	2m41s
nginx-d55b94fd-w2rmv	1/1	Running	0	2m41s

we can get more wider output:

```
root@bakumar2c:/opt# kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE
busybox	1/1	Running	0	70s	10.244.2.4	bakumar4c.mylabserver.com	<none>
nginx-d55b94fd-r25rc	1/1	Running	0	3m14s	10.244.1.4	bakumar3c.mylabserver.com	<none>
nginx-d55b94fd-w2rmv	1/1	Running	0	3m14s	10.244.2.3	bakumar4c.mylabserver.com	<none>

Now we will try to reach one Pod(Ex: Master node-busybox-10.244.2.4) to another Pod(worker node2- nginx-10.244.2.3)

Note:Here 10.244.2.3 is nginx node3 and busybox is master node1.

```
root@bakumar2c:/opt# kubectl exec busybox -- curl 10.244.2.3
```

```
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
   Dload  Upload  Total    Spent    Left  Speed
100  612    100  612    0    0  525k    0 --:--:-- --:--:-- --:--:--  597k
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

It is possible to interact one POD to another POD using virtual network of Kubernetes.(As shown in “Networking in Kubernees ” doc above)

Networking is an important part of understanding the basics of Kubernetes. This lesson provides a high-level overview of what a Kubernetes virtual cluster network looks like. We will also demonstrate how the network functionality connects one pod from another pod over the virtual network.

- Create a deployment with two nginx pods:

```
cat << EOF | kubectl create -f -
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - containerName:
            name: nginx
            image: nginx:1.15.4
            ports:
              - containerPort: 80
  min
  max
```

- Create a service used to serve for routing:

```
cat << EOF | kubectl create -f -
apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
spec:
  selector:
    app: nginx
  ports:
    - port: 80
      targetPort: 80
  min
  max
```

- Get the IP addresses of your pods:

```
kubectl get pods -o wide
```

- Get the IP address of one of the nginx pods, then connect that nginx pod from the busybox pod using the nginx pod's IP address:

```
kubectl exec busybox -- cat /dev/tcp/<ip>/80
```

Kubernetes Architecture and Components

A Kubernetes cluster is made up of multiple individual components running on the various machines that are part of the cluster. In this lesson, we will briefly discuss the major Kubernetes software components and what each of them do. We will also look into how these components are actually running in our cluster currently.

Here are the commands used in this lesson:

- Get a list of system pods running in the cluster:

```
kubectl get pods --all-namespaces
```

- Check the status of the **kubelet** service:

```
sudo systemctl status kubelet
```

Introduction

Building a Cluster

Basic Kubernetes Concepts

Containers and Pods

Clustering and Nodes

Networking in Kubernetes

Kubernetes Architecture and Components

Deploying to Kubernetes

Microservices

Additional Resources

Kubernetes Architecture and Components

Kubernetes includes multiple components that work together to provide the functionality of a Kubernetes cluster.

The control plane components manage and control the cluster:

- **etcd**: Provides distributed, synchronized data storage for the cluster state.
- **kube-apiserver**: Serves the Kubernetes API, the primary interface for the cluster.
- **kube-controller-manager**: Bundles several components into one package.
- **kube-scheduler**: Schedules pods to run on individual nodes.

In addition to the control plane, each node also has:

- **kubelet**: Agent that executes containers on each node.
- **kube-proxy**: Handles network communication between nodes by adding firewall routing rules.

With **kubeadm**, many of these components are run as pods within the cluster itself.

Kubelet

It is an agent run in each Kubernetes cluster node and acts as a middle man between container (Docker) and Kubernetes API.

```
root@bakumar2c:~# systemctl status kubelet
```

```
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset: enabled)
   Drop-In: /etc/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Mon 2019-10-14 06:16:35 UTC; 1h 19min ago
     Docs: https://kubernetes.io/docs/home/
  Main PID: 27212 (kubelet)
    Tasks: 17 (limit: 2318)
   CGroup: /system.slice/kubelet.service
            └─27212 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/
               kubelet.conf --config=/
```

deployments

```
cat <<EOF | kubectl create -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.15.4
        ports:
        - containerPort: 80
EOF
```

root@bakumar2c:~# kubectl get deployments

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx	2	2	2	2	39m
nginx-deployment	2	2	2	2	22s

root@bakumar2c:~# kubectl get deployments -o wide

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
nginx	2	2	2	2	39m	nginx	nginx:1.15.4	app=nginx
nginx-deployment	2	2	2	2	27s	nginx	nginx:1.15.4	app=nginx

Wide description

oot@bakumar2c:~# kubectl describe deployments nginx-deployment

```
Name:          nginx-deployment
Namespace:     default
CreationTimestamp:  Mon, 14 Oct 2019 07:43:11 +0000
Labels:        app=nginx
Annotations:   deployment.kubernetes.io/revision: 1
Selector:      app=nginx
Replicas:      2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType:  RollingUpdate
MinReadySeconds:  0
RollingUpdateStrategy: 25% max unavailable, 25% max surge

Pod Template:
  Labels: app=nginx
  Containers:
    nginx:
      Image:   nginx:1.15.4
      Port:    80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:      <none>
      Volumes:      <none>
  Conditions:
    Type             Status Reason
    ----             -
    Available         True  MinimumReplicasAvailable
    Progressing       True  NewReplicaSetAvailable
    OldReplicaSets:   <none>
    NewReplicaSet:    nginx-deployment-d55b94fd (2/2 replicas created)
  Events:
    Type Reason Age From Message
    ---
    Normal ScalingReplicaSet 9m54s deployment-controller Scaled up replica set nginx-deployment-d55b94fd to 2
root@bakumar2c:~#
```

Introduction

Building a Cluster

Basic Kubernetes Concepts

Containers and Pods

Clustering and Nodes

Networking in Kubernetes

Kubernetes Architecture
and Components

Deploying to Kubernetes

Microservices

Additional Resources

Kubernetes Architecture and Components

Kubernetes includes multiple components that work together to provide the functionality of a Kubernetes cluster.

The control plane components manage and control the cluster:

- **etcd**: Provides distributed, synchronized data storage for the cluster state.
- **kube-apiserver**: Serves the Kubernetes API, the primary interface for the cluster.
- **kube-controller-manager**: Bundles several components into one package.
- **kube-scheduler**: Schedules pods to run on individual nodes.

In addition to the control plane, each node also has:

- **kubelet**: Agent that executes containers on each node.
- **kube-proxy**: Handles network communication between nodes by adding firewall routing rules.

With **kubeadm**, many of these components are run as pods within the cluster itself.

Introduction

Building a Cluster

Basic Kubernetes Concepts

Deploying to Kubernetes

Kubernetes Deployments

Kubernetes Services

Microservices

Additional Resources

Kubernetes Deployments [Next >](#)

Pods are a great way to organize and manage containers, but what if I want to spin up and automate multiple pods?

Deployments are a great way to automate the management of your pods. A deployment allows you to specify a **desired state** for a set of pods. The cluster will then constantly work to maintain that desired state.

For example:

- **Scaling:** With a deployment, you can specify the number of replicas you want, and the deployment will create (or remove) pods to meet that number of replicas.
- **Rolling Updates:** With a deployment, you can change the deployment container image to a new version of the image. The deployment will gradually replace existing containers with the new version.
- **Self-Healing:** If one of the pods in the deployment is accidentally destroyed, the deployment will immediately spin up a new one to replace it.



Introduction

Building a Cluster

Basic Kubernetes Concepts

Deploying to Kubernetes

Kubernetes Deployments

Kubernetes Services

Microservices

Additional Resources

< Prev

Kubernetes Deployments

Let's create a simple deployment!

We'll make a deployment that includes two replicas running basic Nginx containers.

```
cat <<EOF | kubectl create -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
labels:
  app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.15.4
        ports:
        - containerPort: 80
EOF
```

```
kubectl get deployments
```

```
kubectl describe deployment nginx-deployment
```

```
kubectl get pods
```

Deployments are an important tool if you want to take full advantage of the automation capabilities provided by Kubernetes. In this lesson, we will discuss what deployments are and briefly mention some common use cases for Kubernetes deployments. We will also create a simple deployment in our cluster and explain how we can manage it with kubectl. Here are the commands used in this lesson:

- Create a deployment:

```
kubectl create deployment nginx --image nginx:alpine --dry-run -o yaml | kubectl apply -f -
```

- Get a list of deployments:

```
kubectl get deployments
```

- Get a description of a deployment:

```
kubectl describe deployment nginx
```

- Get a list of pods:

```
kubectl get pods
```

You should see the pods created by the deployment.

Kubernetes services

Introduction

Building a Cluster

Basic Kubernetes Concepts

Deploying to Kubernetes

Kubernetes Deployments

Kubernetes Services

Microservices

Additional Resources

Kubernetes Services

Next >

Services are another important component of deploying apps with Kubernetes.

Services allow you to dynamically access a group of replica pods. Replica pods are often being created and destroyed, so what happens to other pods or external entities which need to access those pods?

A Service creates an abstraction layer on top of a set of replica pods. You can access the service rather than accessing the pods directly, so as pods come and go, you get uninterrupted, dynamic access to whatever replicas are up at the time.

Cluster

```
graph LR; S[Service] --> P1[Pod]; S --> P2[Pod]; S --> P3[Pod];
```

Kubernetes services

While deployments provide a great way to automate the management of your pods, you need a way to easily communicate with the dynamic set of replicas managed by a deployment. This is where services come in. In this lesson, we will discuss what services are in Kubernetes, demonstrate how to create a simple service, and explore that service in our own cluster.

Here are the commands used in the demonstration:

- Create a NodePort service on top of your nginx pods:

```
cat << EOF | kubectl create -f -
kind: Service
apiVersion: v1
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      nodePort: 30080
  type: NodePort
EOF
```

- Get a list of services in the cluster:

```
kubectl get svc
```

You should see your service called `nginx-service`.

- Since this is a NodePort service, you should be able to access it using port 30080 on any of your cluster's servers. You can test this with the command:

```
curl -I localhost:30080
```

You should get an HTML response from nginx!

```
root@bakumar2c:~# cat << EOF | kubectl create -f -
```

```
> kind: Service
> apiVersion: v1
> metadata:
>   name: nginx-service
> spec:
>   selector:
>     app: nginx
>   ports:
>     - protocol: TCP
>       port: 80
>       targetPort: 80
>       nodePort: 30080
>   type: NodePort
> EOF
```

service/nginx-service created

```
root@bakumar2c:~# kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	111m
nginx-service	NodePort	10.97.179.188	<none>	80:30080/TCP	14s

```
root@bakumar2c:~# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	111m
nginx-service	NodePort	10.97.179.188	<none>	80:30080/TCP	17s

```
root@bakumar2c:~#
```

```
root@bakumar2c:~# curl localhost:30080
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

```
root@bakumar2c:~#
```

or

```
root@bakumar2c:~# curl 172.31.40.125:30080
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<body>
<h1>Welcome to nginx!</h1>
<
```

Here "172.31.40.125" is the Master node private IP.

Similary using Private IP Addresses of worker nodes ,we can connect to pods/containers.

from workernode2

```
root@bakumar2c:~# curl 172.31.41.71:30080
```

```
<head>
<title>Welcome to nginx!</title>

<h1>Welcome to nginx!</h1>
```

Kuberenetes Microsercvices and Deployments

[Introduction](#)[Building a Cluster](#)[Basic Kubernetes Concepts](#)[Deploying to Kubernetes](#)[Microservices](#)[What are Microservices?](#)[Deploying the
Robot Shop App](#)[Additional Resources](#)

What are Microservices? [Next >](#)

One of the best ways to demonstrate the value of Kubernetes is in managing microservice applications.

Microservices are small, independent services that work together to form a whole application.

Many applications are designed with a **monolithic architecture**, meaning that all parts of the application are combined in one large executable.

Monolith Application

Auth

Customer
DataProduct
Data

Search

Microservice architectures break the application up into several small services.

Microservice Application

Auth

Customer
DataProduct
Data

Search

Search

Product
Data

Search

Search

[Introduction](#)[Building a Cluster](#)[Basic Kubernetes Concepts](#)[Deploying to Kubernetes](#)[Microservices](#)[What are Microservices?](#)[Deploying the
Robot Shop App](#)[Additional Resources](#)[< Prev](#)

What are Microservices?

Here are a few advantages of microservices:

- **Scalability:** Individual microservices are independently scalable. If your search service is under a large amount of load, you can scale that service by itself, without scaling the whole application.
- **Cleaner code:** When services are relatively independent, it is easier to make a change in one area of the application without breaking things in other areas.
- **Reliability:** Problems in one area of the application are less likely to affect other areas.
- **Variety of tools:** Different parts of the application can be built using different tools, languages, and frameworks. This means that the right tool can be used for every job!

Implementing microservices means deploying, scaling, and managing a lot of individual components! Kubernetes is a great tool for accomplishing all of this. In the world of microservices, the benefits of Kubernetes really shine!

Introduction

Building a Cluster

Basic Kubernetes Concepts

Deploying to Kubernetes

Microservices

What are Microservices?

Deploying the Robot Shop App

Additional Resources

Deploying the Robot Shop App

Now we are ready to get hands-on with microservices in Kubernetes. In this lesson, we will deploy a sample microservice application called **Stan's Robot Shop**. This is an open-source sample microservice app made by [Instana](#).

Let's begin by cloning the robot-shop Git repository. This repository contains ready-made YAML files that we can use to quickly and easily install the application.

```
cd ~/
git clone https://github.com/linuxacademy/robot-shop.git
```

Now we can install the app in our cluster, under a namespace called **robot-shop**.

```
kubectl create namespace robot-shop
kubectl -n robot-shop create -f ./robot-shop/K8s/descriptors/
```

Let's check on the pods in the app as they come up!

```
kubectl get pods -n robot-shop -w
```

Once the pods are up, you should be able to access the app in your browser! Use the public IP of one of the nodes in your cluster and port 30080.

```
http://$kube_server_public_ip:30080
```

Kubernetes is a powerful tool for managing and deploying microservice applications. In this lesson, we will deploy a microservice application consisting of multiple value components to our cluster. We will also explore the application itself in order to get a sense of what a microservice application might look like and how it might run in a Kubernetes cluster.

Here are the resources used in the lesson to create and deploy the Stan's Robot Shop application.

- Clone the Git repository

```
cd ~/
git clone https://github.com/linuxacademy/robot-shop.git
```

- Create a namespace and deploy the application objects to the namespace using the deployment descriptors from the Git repository.

```
kubectl create namespace robot-shop
kubectl -n robot-shop create -f https://raw.githubusercontent.com/linuxacademy/robot-shop/master/K8s/descriptors/
```

- Get a list of the application's pods and wait for all of them to finish starting up.

```
kubectl get pods -n robot-shop -w
```

- Once all the pods are up, you can access the application in the browser using the public IP of one of your Kubernetes nodes and port 30080.

```
http://$kube_server_public_ip:30080
```

<http://18.140.237.207:30080/>

Note: here 18.140.237.207 is the public IP address of master node.s

You can see a UI for signup and registration.

Scaling up:

Deploy the Stan's Robot Shop app to the cluster.



1. Clone the Git repo that contains the pre-made descriptors:

```
cd -/  
git clone https://github.com/linuxacademy/robot-shop.git
```

2. Since this application has many components, it is a good idea to create a separate namespace for the app:

```
kubectl create namespace robot-shop
```

3. Deploy the app to the cluster:

```
kubectl -n robot-shop create -f ~/robot-shop/K8s/descriptors/
```

4. Check the status of the application's pods:

```
kubectl get pods -n robot-shop
```

5. You should be able to reach the robot shop app from your browser using the Kube master node's public IP:

```
http://$kube_master_public_ip:30080
```

Scale up the MongoDB deployment to two replicas instead of just one. ^

1. Edit the deployment descriptor:

```
kubectl edit deployment mongodb -n robot-shop
```

2. You should see some YAML describing the deployment object.

- Under `spec:`, look for the line that says `replicas: 1` and change it to `replicas: 2`.
- Save and exit.

3. Check the status of the deployment with:

```
kubectl get deployment mongodb -n robot-shop
```

After a few moments, the number of available replicas should be 2.

Deploying a Simple Service to Kubernetes

Create a deployment for the store-products service with four replicas.

```
cat << EOF | kubectl apply -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: store-products
  labels:
    app: store-products
spec:
  replicas: 4
  selector:
    matchLabels:
      app: store-products
  template:
    metadata:
      labels:
        app: store-products
    spec:
      containers:
        - name: store-products
          image: linuxacademycontent/store-products:1.0.0
          ports:
            - containerPort: 80
EOF
```

Create a store-products service and verify that you can access it from the busybox testing pod.

```
cat << EOF | kubectl apply -f -
kind: Service
apiVersion: v1
metadata:
  name: store-products
spec:
  selector:
    app: store-products
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
EOF
```

```
root@bakumar2c:~# kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3h34m
store-products	ClusterIP	10.110.241.128	<none>	80/TCP	6m46s

```
root@bakumar2c:~# kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx	2	2	2	2	166m
nginx-deployment	2	2	2	2	128m
store-products	4	4	4	4	7m48s

```
root@bakumar2c:~# kubectl exec busybox -- curl -s store-products
```

```
{
  "Products":[
    {
      "Name":"Apple",
      "Price":1000.00,
    },
    {
      "Name":"Banana",
      "Price":5.00,
    },
    {
      "Name":"Orange",
      "Price":1.00,
    },
    {
      "Name":"Pear",
      "Price":0.50,
    }
  ]
}
```

[illegible]

[illegible]

[illegible]
